

# Data Storage and Indexing

Dr Paolo Guagliardo

`dbms-lecturer@ed.ac.uk`



THE UNIVERSITY of EDINBURGH  
**informatics**

Fall 2018

## Files, records and pages

Each table is stored on disk in a **file of records**

**Record**: memory area (sequence of bits) logically divided in **fields**

Each record in a file

- ▶ corresponds to a row of values in the table
- ▶ has the **same number of fields**  
but **not necessarily the same length**
- ▶ has a unique identifier: the record id (**rid**)

Files are organized in **pages**: blocks of memory of fixed size

The **page size** is a parameter of the DBMS

When data is requested for computation

pages must be **fetches from disk** and **loaded in main memory**

# File of records

Supports the following operations:

Insertion of records

Deletion of records

Modification of records

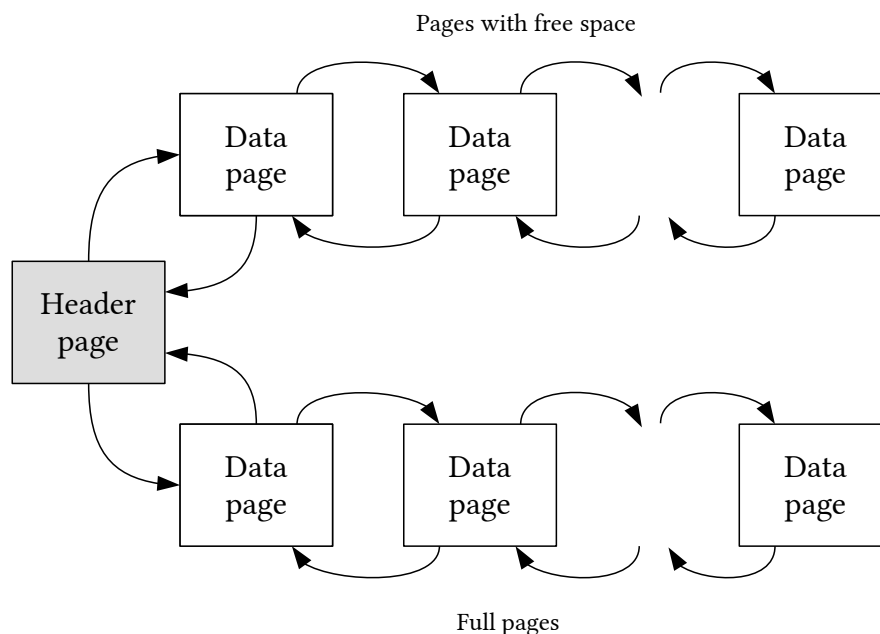
Scan of all records, returned one at a time

Simplest structure: unordered file, called **heap file**

- ▶ records are stored in random order across the pages
- ▶ supports retrieval of a specific record given its rid

Indexed structures allow to efficiently retrieve records that satisfy a given search condition


## Implementing heap files: Linked list of pages



### Disadvantages

- ▶ Almost all pages on free list if records are of variable length
- ▶ Must scan and examine several pages to insert a record

## Implementing heap files: Directory of pages



Free space can be managed by maintaining:

- ▶ a bit per entry  
(free space yes/no)

or

- ▶ a count per entry  
(amount of free space)

## Page formats

A page can be thought of as a collection of slots


- ▶ a record is identified by the **page id** and **slot number**  
so  $rid = (\text{page id}, \text{slot number})$
- ▶ **alternative**: assign unique integer to each record  
and maintain correspondence between rid and (page, slot)

Format of pages depends on:

- ▶ Fixed- vs. variable-length records
- ▶ Support for search, insertion, deletion of records

## Page formats for fixed-length records

### Packed



Records stored in the first  $n$  slots

Records located by offset calculation


Free space contiguous at the end

When a record is deleted,  
the last one is moved to empty slot

Problem if rid contains slot number

## Page formats for fixed-length records

### Unpacked, Bitmap



Bit array tells which slots are free


Records located by offset calculation

Scanning all records requires  
bit array scan + offset calculation

Insertion of record requires  
bit array scan + offset calculation

When a record is deleted,  
corresponding bit is turned off

## Page format for variable-length records



## Page format with directory of slots


- ▶ Most flexible format
- ▶ **Records can be moved without changing rid**
- ▶ Can be used also for fixed-length records
- ▶ Deletion accomplished by setting slot offset to  $-1$
- ▶ Free space must be managed more carefully (the page is not pre-formatted into slots)
- ▶ Only last slot entry can be removed from directory
- ▶ On insertion of record, look for slot entry with offset  $-1$  (if there is none, add new entry to slot directory)

## Fixed-length records

Each field has a **fixed** length (available in the [system catalog](#))

Given the [base address](#)  $B$  of the record,

the address of field  $i$  can be calculated as  $B + \sum_{k=1}^{i-1} L_k$




Direct access to fields, but inefficient storage (especially for nulls)

## Variable-length records


Some of the fields have **variable** length

[Fields delimited by special symbol](#)



Access to fields requires a scan of the record

[Array of field offsets](#)



Direct access to fields ; efficient storage of nulls

## Modifying fields in a record

Potential issues with variable-length records:

- ▶ When field grows, shift subsequent fields to make space
- ▶ If modified record does not fit in the free space on page, move it to another page leaving a **forwarding address** (we must allocate minimum space for every record)
- ▶ If modified record does not fit on any page, split into smaller records across several pages using pointers

Adding fields can cause similar issues in all record formats

## Indexing

### Index

Data structure that organizes data records based on a **search key** (any subset of the fields of a relation)

- ▶ supports efficient retrieval of all data records satisfying a given condition on the search key

Two main indexing strategies

- ▶ **Hashing** (good for conditions based on equality)
- ▶ **Trees** (good for conditions based on ordering)

## Hash-based indexing

Organize records into **buckets**  
based on a **hash function**  $h$  applied to the search key value

For record  $\bar{r}$  and search key  $k$

$$h(\pi_k(\bar{r})) = \text{bucket where } \bar{r} \text{ belongs}$$

Bucket = **primary page** + zero or more **overflow** pages

Example on blackboard

## Tree-based indexing

Records are organised using a **hierarchical tree structure**  
that directs the search from the root to relevant pages

**Non-leaf nodes** contain pointers  $p$  separated by search key values  $v$

$$p_0, v_1, p_1, v_2, p_2, \dots, v_n, p_n$$

For each value  $v_i$

- ▶  $p_{i-1}$  points to a node with values less than  $v_i$
- ▶  $p_i$  points to a node with values greater than or equal to  $v_i$

**Leaf nodes** are pages of data records

### B-tree

Balanced tree: all paths from root to leaves have same length

Example on blackboard